

Backend Developer Take-Home Assignment

Objective

Build a **backend service** that tracks and manages a **custom equal-weighted stock index** comprising the **top 100 US stocks** by daily market capitalization. Each stock must maintain an **equal notional weight**, and the index should update daily to reflect market cap changes and rebalance composition.

Your backend must expose **RESTful APIs** to:

- Construct the index dynamically for any date or date range.
- Retrieve historical performance and compositions.
- Detect and show composition changes over time.
- Export all outputs into **well-formatted Excel files**.

This project assesses your ability to:

- Model real-world financial problems.
 - Build modular services using **FastAPI**.
 - Design efficient **SQL-based** storage with **Redis caching**.
 - Apply backend best practices (APIs, persistence, caching, containerization).
-

Assignment Requirements

1. Data Acquisition (Background Job)

- Fetch daily market capitalization and price data from a public API (e.g., Yahoo Finance, Alpha Vantage, IEX Cloud, Polygon).
- Store the data into **DuckDB** or **SQLite**.
- Evaluate at least **two data sources** and justify your choice.
- Ensure **at least 30 trading days** of history are ingested.
- Data acquisition must run as a **standalone job**, not via API.

2. Data Storage

- Use **DuckDB** or **SQLite**.
- Model tables for:
 - Stock metadata
 - Daily stock price and market cap
 - Index compositions and performance

- Query and transform data purely via **SQL**.

3. Index Construction API

- **POST /build-index**
 - Inputs: start_date, end_date (optional).
 - Construct the equal-weighted index dynamically for the given dates:
 - Select **top 100 stocks** daily by market cap.
 - Assign **equal weights**.
 - Persist compositions and performance.
 - **Cache results in Redis**.
- **Important:** Index building must happen at **API runtime**, not during ingestion. Ensure you understand the principles behind equal weighting, including how notional weights are assigned and rebalanced.

4. Index Retrieval APIs

Endpoint	Function
GET /index-performance?start_date=&end_date=	Return daily returns and cumulative returns (cached).
GET /index-composition?date=	Return 100-stock composition for a given date (cached).
GET /composition-changes?start_date=&end_date=	List days when composition changed, with stocks entered/exited (cached).

- Responses must be **JSON**.

5. Data Export API

- **POST /export-data**
 - Export:
 - Index performance
 - Daily compositions
 - Composition changes
 - Format: **Excel (.xlsx)** with clean headers and numeric formatting.

6. Containerization

- **Dockerize** the entire application.
 - Provide a **docker-compose.yml** bringing up:
 - FastAPI service
 - Local database
 - **Redis instance**
 - Ensure clean networking between services.
-

Tech Stack

- **Python**
 - **FastAPI**
 - **Redis** or similar
 - **DuckDB** or **SQLite**
 - **Polars** and/or **Pandas**
 - **Docker** (with docker-compose)
-

Deliverables

- GitHub repository with:
 - Full source code.
 - **README.md** documenting:
 - Setup instructions (local + Docker).
 - Running data acquisition job.
 - API usage (sample curl/Postman).
 - Database schema overview.
 - Working, clean REST APIs.
 - Working Excel exports.
 - Short write-up suggesting production/scaling improvements.
-

Evaluation Criteria

Category	Focus
Domain Understanding	Correct index logic, rebalancing, change tracking
Backend Engineering	API design, modular structure
Database Design	Schema quality, SQL proficiency
Caching Strategy	Correct and efficient use of Redis
Code Quality	Clarity, modularity, maintainability
Deployment	Working Docker setup
Documentation	Clear and complete